Philosophical Methodology and Levels of Generality

Sarah McGrath
smcgrath@princeton.edu

I. Introduction

Suppose that we're engaged in philosophical theorizing: we're doing ethics, or metaphysics, or epistemology. And we're trying figure out which claims to accept and which to reject. It's a familiar fact that the propositions with which we are concerned in philosophical theorizing range widely with respect of their level of generality. At one end of the spectrum, there are very general principles. For example, in ethics, we find people endorsing extremely general claims like: "One is morally required to act so as to make the world as good as possible." And at the other end of the spectrum, we have claims about the moral status of token actions, such as, "I'm not doing anything wrong by writing this sentence." (And of course, there are many intermediate levels, between very general principles and claims about particular token actions.) Similarly, in metaphysics we find people endorsing extremely general principles about what there is or isn't. For example, consider the claim, "there are no vague objects." (That is, any genuine vagueness must be in our concepts or in the language that we use to talk about the world, as opposed to in the world itself.) And again at the other end of the spectrum, we have claims about the existence of token particulars, such as G.E. Moore's "Here is a hand."

Here are some questions: From the perspective of philosophical methodology, how (if at all) does level of generality *matter*? Is it true that the proper starting point for (say)

theorizing in ethics consists of judgments that have some specific level of generality in common? Is it true (as has also been suggested) that among the judgments that strike us as intuitively plausible, we're more likely to be mistaken at some levels of generality than at others? If so, what follows from that, about how we should proceed? Are there any methodological norms—even highly defeasible norms—about how we should take into account levels of generality when we're engaged in doing philosophy?

Consider, as a possible answer to this last question, the following claim:

> **The General has Priority**: In investigating questions in domain X, we should prioritize very general claims in X.

About ethics, Henry Sidgwick seemed to hold something like this view, both in his practice and in his methodological remarks. Sidgwick thought that we are much more reliable in our judgments about abstract principles than we are in our lower-level judgments, and that our lower-level judgments are "loose, shifting, and mutually contradictory" (ME 338). By contrast, he thought there are very abstract moral axioms that "present themselves as self-evident; as much (e.g.) as the mathematical axiom that 'if equals be added to equals the wholes are equal.'" (ME 383, 507; Hurka p. 110) And in his metaethics, G.E. Moore was basically in agreement with Sidgwick.[1] Now, partly in response to Moore and Sidgwick—against the view that there are very general self-evident principles that are the proper starting points for ethical theorizing—some early-20th century moral philosophers argued that in ethics, something more like the opposite is true. So consider:

> **The Particular has Priority**: In investigating questions in domain X, we should prioritize particular claims in X.

[1] Moore, Sidgwick, as presented by Hurka on pages 120-122.

Here are few examples. The Oxford philosopher E.F. Carritt thought that "…particular judgments about the right act in some particular situation, real or imaginary, are our only data…" (Theory of Morals 31, also 70-1, 84-5, 114-15, 138-39). According to A.C. Ewing, if an intuition concerns general principles, then this "greatly reduces its authority and value" (Morality of Punishment p. 119n).[2]

And of course, there are other views; some might think that claims at an intermediate level of generality are privileged, or endorse a norm that instructs us to take levels of generality into consideration in some more subtle or complicated way.[3] But in what follows, I'm going to try to convince you that it is a mistake to take facts about levels of generality to have any methodological significance at all.

II. The Significance of these Issues for Philosophical Theorizing

But before I get to that I want to say a little bit about one reason why this issue matters. Now it's obvious that where a philosopher comes down on this issue can make a big difference to where she ends up. But I also think that how a philosopher comes down on this issue tends to matter in a specific way.

---

[2] Now, to be clear, these philosophers were not moral particularists in the sense of denying that there are any general principles to be had: they thought that general rules could be formulated, and they thought that such rules are sometimes useful in moral thinking. But, in contrast to Sidgwick and Moore, they thought that our evidence for such principles is always going to be particular moral judgments. (Hurka pp. 122-123).

[3] H.A. Prichard and W.D. Ross disagreed thought, against the generalists, they denied that we can know abstract moral principles just by intuition of the principle itself; they thought that our intuitions about particular cases always plays a role in moral knowledge. But they differed from Caritt and Ewing because they thought that, even about the particular case, what is known is always "implicitly general"; not that an act has a particular moral property, but rather that "an act's having a certain property tends to make it right and will therefore tend to make any other act with that property right." (p. 123).[3]]]

So, start by considering a particularly provocative claim recently made by Scott Soames—in the first volume of his book *The Analytic Tradition in Philosophy*—about Moore's impact on the history of analytic philosophy. Soames suggests that, in a crucial respect, G.E. Moore sets 20th century metaethics off on the wrong foot, dropping the ball by failing to apply, in his ethical writings, the same kind of "bottom up" methodology that he employs in his defenses of common sense in both metaphysics and epistemology. So Soames suggests that there's a deep difference between, on the one hand, Moore's methodology when he's defending common sense against either the revisionary metaphysics of the British idealists or skepticism in epistemology, and, on the other hand, the methodology of *Principia Ethica*. In the papers defending common sense, it's Moore's *opponent* who is employing some very general principle about what knowledge requires, or about the true nature of reality, and then reasoning from these general philosophical principles to radically revisionary conclusions about the extent of our knowledge, or what the world is actually like. And *Moore* is showing how to resist this kind of philosophizing by suggesting, through the example of his practice, that the proper starting points for philosophy are very different. According to Moore, the proper starting points for philosophy are relatively lower-level claims like "Here is a hand" or "I know that this is a pencil." So Moore is reasoning from "the bottom up" and drawing conclusions about the correctness of proposed principles and theories of knowledge on the basis of highly specific knowledge claims.

But Soames reads Moore as having a "flawed conception of justification in ethics," because in ethics, Moore assumes that "ethical justification flows from the general to the particular." Soames continues by saying that for Moore,

> Ethical judgments about particular cases are justified by subsuming them under general moral principles. General principles are justified by appeal to still more general, self-evident moral principles…[4]

But Soames doubts there *are* any unrestricted moral generalizations that have the status Moore says they have: the status of self-evident, foundational moral truth. Soames claims that the fact that many of Moore's contemporaries recognized this played a significant historical role in the rise of emotivism. So the way Soames tells the story, this aspect of Moore's work in metaethics was an unintended stimulus to "the moral skepticism" of Moore's own day.

Soames thinks that Moore should have done the same thing in ethics that he did in his writings on common sense, and that if he had, then he'd have arrived at the more defensible view that "the starting point in ethics consists in our pre-theoretic moral certainties about particular cases and severely restricted generalities" (p.246).

I'm not a Moore scholar, but I think that there's a lot of truth in the story that Soames tells. And I think that there might be a lesson that generalizes beyond this particular case. Soames's discussion of Moore provides an illustration of what might be a more general phenomenon: that you open the door to ending up with *radically revisionary conclusions*, when you start to think that the way to proceed in philosophy is in a top-down (as opposed to bottom up) direction. In other words, I think that Soames' discussion suggests a more general hypothesis: you end up being revisionary if you think that the proper starting point is some very *general* principle or principles that you then consistently apply. And I think that we see this both in the history of philosophy, and

---

[4] "… like the appreciation of beautiful objects and the pleasure of human companionship and interaction are the only things that are intrinsically good, which may be put in the form [For all x, x is good iff x is D]" (p.199, cf. pp. 245-246].

also in contemporary philosophy. And I'll just give three quick examples to illustrate this— one from ethics, one from epistemology, and one from metaphysics. (Of course, it would take a lot more to actually substantiate the thesis in any serious sense.)

An obvious example from ethics is the tradition of revisionary consequentialism, as represented by contemporary philosophers such as Shelly Kagan and Peter Singer. They start off thinking that some very general consequentialist principle is true—say, that we are morally required to perform whichever action available to us that would make the world as good as possible—and then they apply that principle in order to determine the goodness and badness of particular actions. And this leads them to views that are often radically at odds with so-called "common sense morality."

For example, on Kagan's view, almost all of us are acting immorally almost all of the time, because pretty much no matter what we're actually doing, there is some alternative action that would have made a greater contribution to the good. So, it's almost certainly immoral for me to be writing this paper—and for you to be reading it, and so on.

An example from epistemology might be that of Hume. On a traditional interpretation, Hume is both a radical skeptic and also, what Roderick Chisholm called "a *methodist*." (That is, someone who both begins from some very general principle about how we acquire knowledge—in Hume's case, an empiricist principle—and then proceeds to draw conclusions about the extent of our knowledge.) That's a very traditional, textbook reading of Hume: Hume opens the Treatise by announcing a commitment to a very general principle of empiricism, and his skepticism follows from relentlessly and consistently applying that principle.

Finally, here is an example from metaphysics. As I said before, some people find it intuitively compelling that "There are no vague objects." But then it's noticed that lots of things, like, say, mountains, don't seem to have sharp identity conditions. And people have treated this as a reason for concluding that ("strictly speaking") there are no mountains. So again, that would be another case of drawing radically revisionary claim from a very general philosophical claim. Contrast the Moorean move of reasoning from the seemingly evident existence of a mountain to: "There must be something wrong with any principle or theory from which it follows that there *are* no mountains."

So I do think that there *is* this general tendency in philosophy: when you find someone holding a radically revisionary view—about what there is, what we know, or what we're morally required to do—more often than not, they get there by embracing relatively general principles, and then following those principles wherever they lead, undeterred by the fact that the principles turn out to have extremely radical implications.

Of course, there are *exceptions* to the general tendency. You sometimes see philosophers who end up with radically revisionary views who *don't* set out from general principles. (Example: Sometimes people argue for radical skepticism in epistemology by taking as their starting point "We don't know that we're not brains-in vats being undetectably deceived." That's not a general principle about knowledge; rather, it's a claim to the effect that we don't know a certain proposition.)[5] So the idea is just that

[5] Other examples include: Singer himself argues for a radically revisionary conclusion in "Famine, Affluence, and Morality" from a claim about the particular case of the child drowning in the pond. He uses this to argue for a general principle which in turn has very revisionary implications. So it is interesting that one of Singer's most significant papers gets a lot of its argumentative force from a claim about a particular case. Also: arguments for skepticism about moral responsibility start from cases in which someone clearly is not

there's an interesting tendency or correlation here. Of course, actually *substantiating* this hypothesis would require much more than just throwing out a handful of examples, as I have done here.

III. An egalitarian proposal

At this point, let me lay some of my own cards on the table. As a normative matter—if we're interested in how philosophy ought to be done—my view is that the level of generality of the propositions that we're concerned with simply doesn't have *any* deep methodological significance. In other words, level of generality simply doesn't matter in any of the ways that philosophers have sometimes taken it to matter in their practice, their theory, or both.

Here are some theses that I endorse. With respect to any of these theses, you could endorse a version that applies to ethics but not epistemology, or to metaphysics but not ethics, or a generalized version that applies to philosophy in general. I am highly sympathetic to the generalized forms of each of these theses, that is, versions that apply across the board, to philosophical theorizing in general. But just to keep things relatively simple, for now I'll concentrate on the versions that apply to ethics.

So consider the claim DIVERSITY:

> **DIVERSITY**: There is no level of generality that is common to the propositions that make up the proper starting points for theorizing in ethics [epistemology, metaphysics]; rather, the proper starting points include propositions of widely varying levels of generality.

---

morally responsible, and then using such cases to motivate general claims which then (they argue) show that in lots of other cases, people aren't morally responsible.

Why might you think that DIVERSITY is true? Here is what I take to be a strong, prima facie case. First, as a general matter, I think that anything that you *know, prior* to engaging in systematic philosophical theorizing, is something that it's fair game to draw upon once you're engaged in such theorizing. In fact, not only is it fine to draw upon your pre-philosophical knowledge, but it would be a methodological mistake not to do that. (That much is a big theme for many philosophers, including Moore, Quine, Soames, Williamson, and others.) And in the case of morality, I think—maybe optimistically—we know quite *a lot* before we begin to engage in systematic philosophical theorizing. I also think that this pre-philosophical moral knowledge isn't restricted to propositions at some specific level of generality, or even to some limited range of levels.

Some of this pre-philosophical knowledge is knowledge of the status of token events and individuals. (For example, "The 9/11 hijackers acted wrongly in flying planes into the World Trade Center.") Some of this is pre-philosophical knowledge is more general than that, and consists of knowledge of moral claims about the relatively detailed scenarios that philosophers often call "particular cases." And I even think that some of it is knowledge of quite *general* and *unqualified* moral claims. For example, consider the proposition *it is wrong to hurt other people intentionally*. I think that proposition is true, and known to be true by ordinary people.

Here are two points of clarification. First, I just said that among the things that we know, pre-philosophically, are generalizations such as, "It is wrong to hurt other people intentionally." Someone might object: "we know nothing of the kind! After all, there are exceptional cases in which it's *not* wrong to hurt other people intentionally. So we don't

know that it is wrong to hurt other people intentionally, because that generalization isn't *true*."

Of course I agree that there are exceptional cases in which it's not wrong to hurt other people. But this shows that the relevant generalizations are false—and therefore, not known—only if they are interpreted as universal generalizations, as opposed to generic generalizations that tolerate exceptions. And I think that interpreting them as generic generalizations is both the more plausible and the more charitable way of understanding these claims, as they are believed by at least many of those who believe them. (And we should notice that even when the relevant moral claims are understood as generics, they still conflict with some revisionary philosophical theories that have actually been maintained—for example, error theories according to which *nothing* we do actually instantiates *any* moral properties.)[6]

The second point of clarification is this: I've been saying that DIVERSITY is true because we have pre-philosophical *knowledge* at all different levels of generality, and that it is appropriate to draw on anything you *know* when you engage in systematic theorizing. I'm putting the main point here in terms of *knowledge*—and putting it that way might suggest the Willamsonian view that your evidence consists in all and only the things that you know. But you don't need to accept that Williamsonian view to get on board with this kind of case for egalitarianism. First, even you think (against Williamson) that your evidence isn't *restricted* to propositions/facts that you know, still, you might think that anything that you *do* know is part of your evidence, and something that you can

---

[6] Some of Soames' other examples of very general moral claims that are objects of pre-philosophical knowledge: *It is right to keep one's promises*; *it is bad when people suffer*.

and should take into account in doing philosophy. So, if that's right, and if it is true that we have pre-philosophical ethical knowledge at various levels of generality, then this would be enough. Second, even if you don't want to do it in terms of *knowledge*, you might appeal to something weaker, like reasonable belief. So there the claim would be that even before we start doing systematic theorizing, we have reasonable beliefs about morality/right and wrong, and these reasonable beliefs don't have some specific level of generality in common; rather, they are very diverse with respect to the relevant dimension. And that's why DIVERSITY holds.

Along with DIVERSITY, I *also* endorse the following claim:

> **NO PRIVILEGE**: When your ethical views conflict, the respective levels of generality of the conflicting propositions is not a reason for resolving the conflict one way rather than another.[7]

I think that NO PRIVILEGE is also supported by a very straightforward line of reasoning: First, whether you should accept a given claim depends on the *epistemic status* of your judging or believing that it's true—for example, whether your doing so would be reasonable. Second, generally speaking, the epistemic status of a judgment or

---

[7] In an unusual case, the level of generality of a claim could be a reason to reject it. For example, if someone told you that you'd been given a drug that would make you extremely reliable when it came to particular judgments and extremely unreliable when it came to judgments at higher levels of generality, and you knew this. It's a familiar point in epistemology that given the right background story, pretty much any fact can be a reason to believing pretty much anything. So I'm sure you can concoct cases in which, given the set up, in that context, the fact that a judgment has a certain level of generality is a good reason for thinking that it's the judgment that's true. So you might need some kind of qualification to handle that kind of case. I put NO PRIVILEGE this baldly, because that does reflect the spirit of what I think is true. Generally speaking, in a case in which you become aware of a conflict between two moral judgments that you've held up until then, the fact that one of the two judgments has a certain level of generality while the other has some different level of generality is not itself a reason—even a prima facie, defeasible reason—for resolving the conflict in one way rather than another.

belief is not determined by, or even significantly correlated with, its level of generality. So, in a case involving conflicting judgments, the questions to ask will be about the comparative epistemic standing of the judgments in question. For example, is either judgment among the things that you know to be true? If not, which of the conflicting judgments is it more reasonable to think is true? It would be strange to think that much progress could be made in answering questions of this sort by reflecting on the levels of generality of the judgments at issue.

The kinds of things that have epistemic statuses are (at least in the first instance) *token* judgments and beliefs. And a token judgment is something that is made by a particular person at a particular time. The epistemic status of a person's token judgments is typically going to be quite sensitive to facts about how she is embedded in the world at the time she makes them. (Different people differ in what evidence that they have; the same person differs in the evidence she has at different times.)

But the level of generality of a token judgment is something that it shares with any other token judgment of the same type—that is, any other token judgment with the same propositional content—regardless of who makes it or how she's embedded in the world at the time that she makes it. So token judgments of the same type will not differ in their level of generality, but they *will* often differ in their epistemic status. This suggests that epistemic status is largely independent of level of generality.

I'm going to use the term "egalitarianism" for the conjunction of these theses, and for the general outlook or idea behind them: namely, that considerations having to do with levels of generality don't have any deep methodological significance.

IV. Challenges to Egalitarianism from Top-Down Ethical Theory

I think that the very straightforward kinds of considerations that I just mentioned amount to a strong prima facie case for egalitarianism. But this way of thinking about things can be, and has been, challenged. And there are any number of different challenges that might be made. I want to look at one kind of challenge that has at times loomed large in ethics, including in recent decades. The challenge is sometimes raised by consequentialists, on the grounds that egalitarianism encourages us to see relatively low-level, common sense moral judgments as giving us good reason to reject more general and abstract consequentialist principles, principles that are alleged (by, for example, Moore, and others) to enjoy something like the status of self-evident axioms. In recent decades, this issue has often come up in discussions of the method of reflective equilibrium, because the egalitarian picture that I endorse is a central commitment of the method. As Tim Scanlon puts it in an exposition and defense of the method: "What the method of reflective equilibrium prescribes is, so to speak, a level playing field of intuitive justification on which principles and judgments of all levels of generality must compete for our allegiance" (2002: 151). (I should say immediately that this idea isn't the exclusive property of the reflective equilibrium theorist. I myself don't think that the method of reflective equilibrium *is* the correct account of justification or inquiry in ethics or anywhere else. But I do think that the reflective equilibrium theorist is right on this point.)

Many consequentlialists have complained about this aspect of the method. So, for example, here is Singer, in the course of criticizing Rawls' view that some of our low-level particular moral judgments count strongly against utilitarianism:

> Why should we not make the opposite assumption, that all the particular moral judgments that we intuitively make are likely to derive from discarded religious systems, from warped views of sex and bodily functions, or from customs necessary for the survival of the group in social and economic circumstances that now lie in the distant past? In which case, it would be best to forget all about our particular moral judgments, and start again from as near as we can get to self-evident moral axioms (p.516)

So I take it that if anything like what's being suggested here is correct, then DIVERSITY and NO PRIVILEGE are incorrect. (As an aside: in the paper from which this passage is taken, Singer says that our moral judgments about particular cases are unreliable because they are more likely to be due to cultural conditioning.[8] In a more recent paper, "Ethics and Intuitions", Singer makes the same claim—our moral judgments about particular cases are unreliable—but there he attributes this to the role that biological and evolutionary programming play in influencing such judgments.)

Other people have also had this idea, that our more particular, lower level moral judgments are especially susceptible to bias, and thus as a class are less reliable than

---

[8] "We have all been making moral judgments about particular cases for many years before we begin moral philosophy. Particular views have been inculcated into us by parents, teachers and society from childhood. ….These judgments sink deep, and become habitual. By contrast, when we read Sidgwick for the first time we are suddenly called upon to decide whether certain fundamental moral principles, which we may never have explicitly thought about before, are self-evident. If it is then pointed out to us that this fundamental moral principle is incompatible with some of the particular moral judgments we are accustomed to making, and that therefore we must either reject the fundamental principle, or else abandon our particular judgments, surely the odds are stacked against the fundamental principle. Most of us are familiar with lingering guilt feelings that occur when we do something that we are quite certain is right, but which we once thought to be wrong. These feelings make us reluctant to abandon particular moral views we hold, but they in no way justify these views" (pp.516-517).

more abstract and general judgments. For example, Michael Huemer claims, "level of generality matters because intuitions of different levels of generality differ in their susceptibility to various kinds of error." According to Huemer, we should give priority to "abstract theoretical intuitions" over both "mid-level" and "concrete intuitions", because abstract theoretical intuitions are less susceptible to being influenced by emotional, cultural and biologically generated biases than the others.

So Singer suggests that it might very well make sense to set aside (or as he puts it "forget all about") our particular moral judgments and simply decide directly which general principles are true. Now there are two questions that we might ask in response to suggestions of this kind. The first is: "Do we really have good evidence that our particular moral judgments are unreliable to the point that it would make sense to do this?" (And for what it's worth, I think that the evidence for thinking that we're unreliable when it comes to making particular moral judgments is often exaggerated.)

But I actually want to set that question aside, and consider a different question about the line that Singer is pushing here. Namely: suppose that we did get hard evidence that we're unreliable when it comes to making particular moral judgments. Where exactly would that leave us? Would the reasonable response to that be to employ the kind of anti-egalitarian methodology that Singer is advocating? I don't think so. Let me explain why.

First, we should bear in mind that the kind of reflective equilibrium theorist that Singer is criticizing doesn't think that you get to rely on just any judgment about a particular case that you might make. She thinks that the relevant judgments are so-called "considered judgments", where your considered judgments are judgments that seem

clearly true to you under conditions that are generally hospitable to making good judgments. (For example, you're not distracted or upset, or anything like that.)

So let's imagine getting evidence that we're unreliable even with respect to our considered judgments about cases. For example, suppose we ask an oracle. Now, before meeting the oracle, there are various general moral principles that seem clearly correct to us, and also lots of particular moral judgments that seem clearly correct to us. And suppose that we only get to ask the oracle one question. So we ask the oracle: "how good are we at making particular moral judgments?" And the oracle says: "You're TERRIBLE! VERY UNRELIABLE! Even your considered judgments—that is, even the judgments that seem clearly true to you, when you're in good circumstances—are unreliable when they are about particular cases."

Let's assume that the testimony of this oracle would be compelling evidence that we're unreliable when it comes to making particular judgments. So the question is: how would it be reasonable to respond to that information?

It seems like Singer thinks the response should be: well, we still have a perfectly good way of conducting moral inquiry. We'll just ignore or set aside what seems true to us, when we're making particular moral judgments. And we'll just decide directly which general moral principles are true, based on which seem true.

The picture here seems to be that we're simultaneously *un*reliable when it comes to making particular moral judgments, but reliable *enough* in our judgments about general principles to have justified confidence in those judgments. (And again, we manage to attain that level of reliability by studiously ignoring the particular moral

judgments that we're naturally disposed to make in virtue of our upbringing, judgments that would otherwise corrupt our assessment of the principles.)

That's a perfectly consistent picture. But I think that it's also somewhat implausible. And there are at least two reasons why. First, it's is implausible because it rests on an implausible picture of our psychology. The picture of our psychology is one on which we might be in a position to confidently endorse a general principle, or choose one general principle over another, in a way that's completely uninfluenced by our sense of the intuitive plausibility of its concrete implications. But it's natural think that you wouldn't even be in a position to fully understand the content of a general moral principle unless you understood at least some of its concrete implications[9], and that once you understand those concrete implications, your sense of their intuitive plausibility will inevitably influence your assessment of the general principle that implies them. So you might think that the procedure, "just figure out which general moral principles to endorse, without being influenced by your sense of the intuitive plausibility of any of the concrete implications," is psychologically unrealistic.

But a more serious worry about the picture is epistemological. Consider the moment when the oracle tells you: "I've got bad news: you're unreliable even when it comes to those particular moral judgments that seem clearly true to you." Getting that information should, I think, significantly undermine your confidence that you're *reliable* when it comes to determining which general moral principles are true, based on which principles seem clearly true to you. It's not as though the oracle tells you: "Well, I've got

[9] William James once said: "No one sees further into a generalization than his own knowledge of the details extends." But you don't have to go that far, in order to think that there's something fishy about the Singer picture.

some good news and some bad news. The bad news is that you're unreliable in judging the truth of lower level moral claims based on which ones seem clearly true to you. But, luckily, there's the Good News: you are reliable when it comes to judging the truth of general moral principles, when you go by which ones seem clearly true to you." You're not getting that kind of information. (And this corresponds to the fact that Singer isn't giving you actual evidence that you're reliable when it comes to choosing general principles; rather, he's suggesting some reasons why you might be *un*reliable with respect to lower level judgments.[10])

So part of my thought here is that even if there's some kind of defeasible presumption that when a general moral principle seems clearly true to me, I'm entitled to endorse it as true, that defeasible presumption would not survive finding out the following thing: that when it comes to lower level moral claims, *what seems clearly true* to me is *not* a good guide at all to what is true.

What this suggests, I think, is the following: the closest possible world (perhaps the actual world) in which we're unreliable even with respect to those lower level moral judgments that seem most clearly true to us is not a possible world in which we nevertheless manage to be reliable when it comes to judging which general moral principles are true. Instead, it's a possible world in which we shouldn't have confidence in the moral judgments that we find ourselves disposed to make at *any* level of generality. Skepticism about our ability to make particular moral judgments with a fair degree of accuracy leads, I think, to skepticism about our ability to make accurate moral judgments,

---

[10] I'll consider below the implications of getting comparative evidence to the effect that we are at least somewhat more reliable with respect to higher level judgments than lower level judgments.

regardless of their level of generality. And of course skepticism about our ability to make accurate moral judgments is not a view according to which we should give priority to some moral judgments over others in virtue of their level of generality. It's a view according to which we should refrain from relying upon moral judgments altogether, to the extent that we're able to do this.

I just addressed is an objection that says: our intuitive judgments at the lower levels of generality are unreliable (say, because they're hopelessly biased), and so we should judge the truth of more abstract, general claims directly. Part of my objection to this proposal is that, if we really were hopeless about judgments at the particular level, we would probably also be hopeless about judgments the abstract, general level. (Or at least, it wouldn't be reasonable for us to think that we are reliable when it comes to such judgments.) Now I want to consider a related objection to egalitarianism. This objection starts by supposing that what we have is a certain kind of *comparative* evidence to the effect that our lower-level moral judgments are somewhat more susceptible to distorting influences than our judgments at higher levels of generality (or vice versa). (Singer's claim is that you have reason to think that your particular moral judgments are untrustworthy. Here, I'm imagining someone who says: it's not *just* that you've got reason to think that you're particular judgments are unreliable. You've actually got hard evidence in favor of the *greater reliability* of your higher-level judgments.)

Again, I don't believe that we actually have that evidence. But I think it's worth considering: what, if anything, would follow if we did? The main point I want to make here is that it is not obvious what anti-egalitarian norms such evidence would motivate, or how to argue from such evidence to the denial of NO PRIVILEGE.

Now of course, we could imagine various extreme hypothetical scenarios. For example, scenarios in which the oracle tells us something like this: "In the future, when you make a judgment about a general principle, you'll always be right, or you'll *almost* always be right. And on the other hand, [the oracle continues] your judgments about particular cases will be so unreliable that you might as well flip a coin." So if we actually got evidence like that, then I agree that we should give up on egalitarianism. But I take it that realistically, we're not going to get evidence like that. (And I also take it that nothing interesting about methodology follows from how it would be reasonable for us to respond to evidence like that in a possible world in which we actually get it.)

Realistically, if we did get hard evidence that bears on the comparative question, it would be more like this: it turns out that our judgments at some level of generality are, as a class, ***somewhat*** more susceptible to distorting influences than our judgments at some other levels of generality. I think that *wouldn't* amount to a decisive reason to think that, say, NO PRIVILEGE must be wrong.

Consider an analogy. Suppose that we had compelling evidence that there is a statistically significant correlation between mathematical ability and gender, such that men are somewhat more likely than women to be above average in mathematical ability. We can imagine that this evidence is strong enough to justify a belief in a statistical generalization to that effect, although it's perfectly consistent with the obvious fact that many women are above average when it comes to math, and many men are below average. So: *that* kind of statistical information would *not* provide a sound basis for judging that any particular man is better than any particular woman with respect to math. In fact, given that many women are above average in mathematical ability, and given that

many men are below average, that kind of statistical information wouldn't justify the judgment that a particular man is superior to a particular woman even if you knew nothing about the two people being compared other than their genders.[11] But even more clearly: statistical information of this kind wouldn't justify judging that a particular man has superior mathematical ability to a particular woman in a case in which you have information about the actual mathematical abilities of the two people being compared.

I think that something similar is true when we think of conflicts between judgments at different levels of generality. When faced with a specific conflict between a general moral principle and a particular moral judgment, it would be a mistake to use purely statistical information of an analogous kind as a basis of resolving the conflict in one way rather than the other. After all, everyone will agree that there are countless moral principles that it would be unreasonable for us to accept. And anyone who isn't a moral skeptic will presumably agree that there are many lower-level moral judgments that it's reasonable for us to accept. I assume that any statistical information that might emerge concerning the relative likelihood of bias for judgments at different levels of generality would be consistent with that much. So in those cases in which it IS reasonable for us to endorse a specific moral principle over an incompatible lower-level judgment, it's not the level of generality of the propositions that makes this the reasonable course. And I think that, although this is true in general, it's especially clear in any case in which you have specific information about the propositions at issue *beyond* their level of

[11] Although in the absence of any other information, such statistical information might justify betting that the man had superior ability, if you were forced to place an otherwise blind bet on whose ability was superior. But the crucial point is that even in those circumstances you would not be justified in believing the relevant proposition; rather, agnosticism would be in order.

generality: information that bears on their belief worthiness. But given that you grasp the contents of the propositions being compared, you typically will have at least some information that bears on their belief worthiness. (It's not as though you're in the position of placing a "blind bet" about which proposition is more likely to be true, given only information about their respective levels of generality.)

To recap: I have just argued that, even if we had a certain kind of comparative evidence that our judgments at a certain level of generality were more susceptible to bias, that wouldn't amount to a decisive reason to think that NO PRIVILEGE must be wrong. We can think of it as motivating a certain kind of challenge. Often you find philosophers talking about this issue and suggesting that, for this or that reason, judgments at such-and-such a level of generality are likely to be less trustworthy, as a class, than judgments of some other level of generality, as a class. And it's clear that they think that something significant follows about how we should be doing philosophy. But rarely do you find them actually proposing some actual methodological norm, a norm that makes reference to levels of generality. Singer is a possible exception, because in the passage on your handout, it seems like he's pretty close to saying: our particular moral judgments are so unreliable, that we should just set all of them aside. That's a very radical proposal. And I don't think that it's especially plausible. (And I'm not sure that is or ever was his considered view.) But at least it's an actual, determinate norm. If you don't go for something like that, then it's at least not obvious what norms for philosophical inquiry would be justified by evidence of the relevant kind.

I've just argued that the fact that a given moral judgment belongs to a certain class of moral judgments in virtue of having the level of generality that it does is not itself a reason to discount that judgment relative to others, even if members of that class are, as a statistical matter, somewhat more susceptible to the influence of various kinds of bias. The last thing I want to do in this section is consider a different approach that a philosopher who holds that we tend to give too much weight to our lower-level judgments in the context of moral inquiry might take.

Such a philosopher might attempt to undermine your lower-level judgments in a more piecemeal fashion, by arguing that specific moral judgments either: align with our self-interest, or are judgments that it would be painful for us to give up or even seriously question, or are ones that we can be expected to hold given facts about cultural conditioning or our evolutionary history, or for some other reason.

Now I think that this approach is at least more promising than any approach that attempts to establish, on general grounds, that judgments of a certain level of generality are epistemologically or methodologically fundamental, or that there are norms for conducting moral inquiry that would prioritize some moral judgments over others in virtue of their level of generality. But we shouldn't underestimate the burdens that face this more piecemeal approach.

For example, even if you know, of some substantive moral view that you hold, that you have extremely strong psychological incentives for holding it, and that these psychological incentives make it likely or even guarantee that you would hold that same view even in the absence of good reasons for thinking that it is true, it doesn't follow that

you should set aside or discount that judgment in the context of moral inquiry. To illustrate the point: suppose a person of color makes the substantive moral judgment

> *A person of color should not receive lesser consideration in virtue of being a person of color.*

She might make this judgment while knowing about herself that she is heavily invested in the truth of this claim, and that she has tremendous psychological incentives to believe it. Indeed, we can imagine that she is informed by a perfectly reliable oracle of the following fact: her psychology is such that she would endorse this judgment *regardless of whether she has good reason to think that it is true.* But even if she knows all of this, it doesn't follow that she should set aside or discount this judgment when she engages in moral inquiry. For all that has been stipulated about the case so far, she might have compelling reasons to believe the relevant moral claim. Indeed, for all that's been stipulated about the case, it might be among the things she knows to be true. But if it's among the things that she knows to be true, then it would be a mistake for her to set it aside when she engages in moral inquiry: here as elsewhere, she should take into account any knowledge that is potentially relevant to the questions she is attempting to answer. And this is so even if she would still hold the belief in the counterfactual scenario in which she did not know it or have good reasons to think that it is true.

The argument that she need not set aside her belief runs as follows: even if her psychology guarantees that she would believe the proposition (and she knows that), that is consistent with her knowing that the proposition is true. But there is no genuine norm that requires you to set aside relevant information that you know to be true. Therefore, there is no genuine norm that would require her to aside the belief simply because her psychology guarantees that she holds it.

V. Concluding Remarks

The last thing I want to do is to say how I differ from a number of philosophers with whom I agree about some important things. So first, there's a group of philosophers that includes people like Roderick Chisholm, Moore in his writings defending common sense (at least as he's often interpreted), and perhaps Soames. And these are philosophers who reject the idea that the proper starting point for philosophy consists in very general, abstract principles, and that philosophical theorizing consists of drawing out the implications of those principles.

I agree with them about that crucial point. But there also seems to be a suggestion in their work that there's some other privileged level of generality, or some specific level of generality that's shared by the judgments that constitute the proper starting point for philosophical inquiry. And that suggestion would be incompatible with the principle DIVERSITY. So to the extent that that's the case, these philosophers are from my perspective insufficiently egalitarian.

This is maybe most explicit in Chisholm. In his classic discussion of "The Problem of the Criterion", Chisholm distinguishes between different approaches to the problem, which correspond to different ways of doing philosophy. First, you can proceed as a "methodist": that's Chisholm's term for someone who thinks that the proper starting point for inquiry consists in some general principle or general method, the correctness of which we have antecedent knowledge. Alternatively, you can proceed as a "particularist": that's someone who thinks that the proper starting point for philosophical

theorizing consists of our judgments about particular cases. (So notice that 'particularism' in Chisholm's sense is very different from 'particularism' in, say, Jonathan Dancy's sense, or the way in which the word is most frequently used in contemporary meta-ethics. In Dancy's sense, 'particularism' is the view that there aren't any true general principles to be had in a given domain. Whereas Chisholm thought that there are general principles to be had, but the way in which we discover those general principles is by theorizing from a starting point that consists of our judgments about cases.)

So for Chisholm, no less than for the methodist, there *is* a privileged level of generality. It's just that the methodist is badly mistaken about what that level of generality is. And it's also natural to read Moore this way, in his writing defending common sense. In fact, Chisholm explicitly claims that Moore is a particularist in his sense. (When he surveys the history of philosophy with this distinction in mind, he says: "G.E. Moore is an excellent example of a particularist in my sense.")

Similarly, although I don't want to flat out attribute this to him as his considered view, some of the things that Soames says in his discussion of Moore's metaethics at least strongly suggest the view that the correct starting points for ethical theorizing do have some level of generality in common. For example, when Soames describes the procedure that he thinks that Moore should have employed in ethics, Soames offers a number of examples of what he takes to be true claims about morality. And each of these claims is a generalization, although they're highly restricted generalizations. (Here are two of Soames's examples, to give you the flavor: "Any man who tortures children to death solely for the pleasure of watching them suffer and die is a bad man"; and "Any

state of affairs in which every sentient being suffers alone in pain with no relief of any kind, followed by death, is bad.") Soames says, "Restricted generalizations like these are the platitudes that constitute our starting points in ethics" (p.200). So here again, the suggestion seems to be that the proper starting points in ethics do have some specific level of generality. And to the extent that that's what's being suggested, I disagree.

To summarize: I agree with this group of thinkers that a lot of mischief ensues when philosophers make a certain methodological error, namely, that of thinking that the way to proceed in philosophy is by reasoning in a top-down manner, drawing conclusions from general principles. But I think that once we correctly reject the idea that general principles are a privileged level, we shouldn't fall into thinking that there's some lower level that has methodological priority.

(Of course, to the extent that we think that a lot of philosophy in the tradition has gone wrong in neglecting or discounting our lower level judgments—that the tradition has a certain "bias against the particular"—it might make sense to compensate by emphasizing the importance (as Soames does) of bringing to bear especially secure lower level judgments when we assess theories and general principles. But I would resist the idea that that there is any sense in which the lower level judgments have some kind of methodological significance in virtue of their level of generality as opposed to their epistemic status. And again, their epistemic status is something that they will share with judgments of different levels of generality.)

So on this cluster of issues, as I said, I agree with the reflective equilibrium theorist: level of generality does not itself have any deep methodological significance.

And, as I said before, I don't think that that idea is the exclusive property of the reflective equilibrium theorist. You can endorse it even if you don't endorse the method of RE.

However, from my perspective, there's a crucial part of the story that the reflective equilibrium theorist is not in a good position to provide. In particular, I think that the reflective equilibrium theorist isn't in a good position to explain *why* a claim like DIVERSITY is correct.

Let me explain what I have in mind. On my view, a crucial part of the story about why DIVERSITY is true is that in ethics—as well as in epistemology in metaphysics—our prephilosophical knowledge of the subject matter isn't restricted to propositions that have some level of generality in common. Rather, it includes knowledge of propositions of all different levels of generality. (And to the extent that I started to think that all of our pre-philosophical knowledge in some domain was at some specific level of generality, then I would start to question DIVERSITY as it applied to that domain.) The story that I want to tell presupposes that we have some full-fledged knowledge of the domain—knowledge about right and wrong, knowledge about what we know, or knowledge about what exists—prior to engaging in philosophical reflection or theorizing. On my view, that's a crucial part of the story.

But I think that that crucial part of the story is not something that the typical proponent of the method of reflective equilibrium is in a good position to tell. If you look at the standard characterizations of the method by its proponents, and see what they say when they address the question: "Which judgments should we start from?"—You won't find them saying things like "Start from what you know" or even something weaker, like "start from the things that it's reasonable for you to think are true at that time." Rather,

they'll say something like: "The correct starting point consists of your considered judgments about the subject matter", where 'considered judgments' are characterized in non-normative terms. (For example, in Rawls' classic discussion, your considered judgments are judgments of which you're confident, are judgments that are stable over time, are ones that you endorse when you're not upset or frightened, and so on.) And I think that one reason why reflective equilibrium theorists insist on characterizing the starting point in non-normative terms is the following: lots of them think of the method as an account of justification in the moral domain, so that your justification for holding a moral view in some sense arises out of applying the method. (For example, in his most recent book, Scanlon says: "the only way we have of establishing the truth of normative judgments is through direct, piecemeal application of the method of reflective equilibrium.")

If that's your picture—that knowledge of the subject matter arises out of employing the method itself—then obviously, that's going to push you away from characterizing the proper starting point as "Start from what you already know about the subject matter!"

So I think that, although the RE theorist will correctly endorse DIVERSITY, he won't be in a good position to give a good explanation of why it's true. His explanation will run something like this. "Look, it turns out that, as a matter of fact, we do have considered judgments at all of these different levels of generality." And then there's this extra norm: "Don't discriminate for or against any considered judgment simply in virtue of its level of generality." But that seems unsatisfying in the face of Singer-type concerns to the effect that: "Well, maybe you should discriminate against some of your considered

judgments. In particular, the mere fact that you've confidently held such-and-such judgments for a long time, and they seem true to you even when you're not upset or frightened, etc. isn't sufficient to show that it's a proper starting point."

The correct explanation of why DIVERSITY is true, I think, is that our pre-philosophical knowledge is itself diverse with respect to level of generality.[12]

[12] Alternatively, if you're suspicious of appeals to knowledge, the story might go like this. When we engage in philosophical inquiry, we should take into account anything that it's reasonable for us to believe about the subject matter, at any given moment in time. And typically, at the outset of philosophical inquiry, there are things that it's reasonable for us to believe about the subject matter, and those things are diverse with respect to their level of generality. Either of those explanations, I think, would be superior to the explanation that the Reflective Equilibrium will offer of why a claim like DIVERSITY holds.