

ALLOWING THE WRONG, TO BRING ABOUT THE GOOD*David Faraci*

Anne, a wealthy business owner, offers Bob, who is quite poor, full-time employment for which he will receive compensation insufficient to meet his basic needs (food, shelter, etc.). Anne could offer Bob better compensation without lowering the net benefit of their interaction or eliminating her profits. Indeed, Bob's compensation represents only a small fraction of the profits his work generates for Anne. Nevertheless, Bob correctly believes that the position on offer is better for him than any feasible alternative—including alternatives he would have if Anne's business did not exist. He therefore consents to work for Anne. No third parties are harmed by their interaction.

To many, this will look like a case of *wage exploitation*: Anne seems to be *taking advantage* of Bob in paying him such a low wage.¹ It may seem that Anne acts wrongly in doing this, even though she had no prior obligation to benefit Bob, both she and Bob consent to the interaction, it makes them both better off, and there are no negative externalities.² Call this view of the case the “Reasonable View.”

In Part I of this paper, I argue that the Reasonable View is indeed reasonable. I do not argue that it is *true*, only that there are perfectly sensible, intuitive ethical positions that vindicate it. The Reasonable View requires such defense in light of the “nonworseness claim,” according to which a consensual, mutually beneficial transaction cannot be “morally worse” than its absence (Wertheimer 1999, 289–93, though note that Wertheimer himself rejects the claim). The apparent upshot for our

¹ For an overview of accounts of wage exploitation (with respect to sweatshops), see Snyder (2010).

² If “exploitation” is a thick moral term (with a negative valence) then saying that Anne acts wrongfully is redundant. I take no position on this here. For arguments for the neutrality of ‘exploitation’, see Valdman (2009).

example is that if it is permissible for Anne not to hire Bob, it cannot be impermissible for her to hire him, even with such low compensation. This may be true: for instance, if welfare and autonomy are all that matters morally. But the nonworseness claim's proponents hold more strongly that the nonworseness claim should be a point of nearly universal agreement, that *any* reasonable moral theory should endorse it. I argue that this is false.³

As will become clear, it can be tempting given my arguments to conclude that only *deontologists* can embrace the Reasonable View.⁴ Indeed, it has been argued that the Reasonable View should be rejected because it faces “the paradox of deontology.” I conclude Part I with a principled response to this objection, on behalf of deontologists. In Part II, I develop a toy consequentialist theory consistent with the Reasonable View. This serves as a further response to the “paradox of deontology” objection, further illuminates the intuitive roots of the Reasonable View, and demonstrates that even consequentialists can endorse it.

Throughout, I assume that the wrongfulness of Anne's action stems from its *unfairness*.⁵ This assumption is dialectically useful for a number of reasons, independently of any particular view of what fairness consists in. First, defenders of interactions like Anne and Bob's are typically, sometimes even explicitly, concerned primarily or exclusively with the values of *welfare* and *autonomy* (e.g., Powell and Zwolinski 2012). *Fairness* is a clear example of something that many take to be valuable, independently of its impact on welfare or autonomy.

³ My argument can be read as an expansion upon Wertheimer's own reasons for rejecting the nonworseness claim. Other critics of the nonworseness claim include Arneson (2013), Bailey (2010), Barnes (2013) and Malmqvist (2015).

⁴ For our purposes, “deontology” is just non-consequentialism.

⁵ Proponents of the view that wage exploitation is wrongful because it is unfair include Andrews (1852), Meyers (2004), Valdman (2009) and Wertheimer (1999). Perhaps the most popular alternative theory is that wage exploitation is wrongful because it is *disrespectful*. See, e.g., Arnold and Bowie (2003) and Snyder (2008). Some accounts seem to appeal to both unfairness and disrespect, e.g., Sample (2003).

Second, it is not clear that our duties to act fairly are duties *to* the people we might treat unfairly. We may object to Anne's getting *more* than is fair whether or not Bob consents to his getting *less*. This is important because it allows us to bypass the worry that even if there is something unfair about what Anne does, Bob can be taken to have waived her duty to him when he consents to work for her, for it may not be his to waive.⁶

Finally, aside from appeals to the nonworseness claim, objections to the view that wage exploitation is unfair are largely empirical or pragmatic—e.g., critics argue that real-world employers typically don't get an unfair portion of the profits generated by their employees.⁷ My arguments show that those attracted to the possibility of wrongfully unfair wage exploitation should feel no pressure from the nonworseness claim, and that anyone who wishes to argue *against* the possibility of wrongfully unfair wage exploitation will need a new principled argument against the moral importance of fairness (at least as compared to other relevant values, and at least in this context). Instead, I suggest, defenders of wage exploitation should be satisfied with a defense that is contingent on the empirical facts.

Part I. The Reasonableness of the Reasonable View

Consider two passages, the first from Wertheimer, introducing the nonworseness claim, the second from Matt Zwolinski, one of the nonworseness claims' champions:

⁶ Powell and Zwolinski (2012) levy this objection against Snyder's (2008) claim that employers like Anne fail to meet their duty of beneficence. Since this would be a duty Anne has *to Bob*, the objection arguably has more weight against Snyder than against the fairness view.

⁷ For example, Powell and Zwolinski (2012) argue that employers don't get much more from the interaction than other employers do from less apparently exploitative arrangements. I stipulated, by contrast, that Bob receives only a small portion of the profits he generates for Anne. Other views of the wrongness of wage exploitation are arguably subject to principled objections independent of the nonworseness claim (see previous note).

Suppose that we criticize A for engaging in an unfair mutually advantageous transaction with B. A defends himself by appealing to what I shall call the *nonworseness claim*: “Look, you grant that I have a right not to transact with B. If B is better off if I engage in an unfair transaction with B than if I do not transact with B at all, it certainly can’t be *morally worse* for me to engage in such a transaction with B. Given that I have a right not to transact with B and that transacting with B is not worse than not transacting with B, it can’t be seriously wrong for me to engage in an unfair transaction with B.” (Wertheimer 1999, 289)

[The nonworseness claim] holds that in cases where A has a right not to transact with B, and where transacting with B is not worse for B than not transacting with B at all, then it cannot be seriously wrong for A to engage in this transaction, even if its terms are judged to be unfair by some external standard . . . From a consequentialist moral framework, the nonworseness claim seems obviously true. (Zwolinski 2008, 357)⁸

Following these lines of thought, the champion of the Reasonable View charges that Anne’s hiring Bob is wrongful because it is unfair. Anne replies that since Bob is better off than before, and she didn’t have to benefit him at all, her benefitting him can’t be wrongful.

On its face, this is a rather strange reply. The obvious rejoinder is that the presence of a benefit isn’t what’s at issue; *unfairness* is. Zwolinski’s contention that “the nonworseness claim seems obviously true” for consequentialists further highlights this oddness. After all, Wertheimer and Zwolinski both *just told us* what moral value Anne’s exploiting Bob might negatively impact: fairness. Nothing prevents consequentialists from acknowledging the value of fairness. And so nothing prevents consequentialists from claiming that even though “transacting with B is not worse **for B** than not transacting with B at all” (bolding mine), it is nevertheless worse *all things considered* because of its unfairness.

As an analogy, suppose someone were to propose a variant of the nonworseness claim that lacks the consent condition—i.e., claim that a morally optional interaction cannot be wrongful just so long as all parties are made better off, even if this is achieved through coercion. It would be baldly

⁸ This passage is from a paper about price-gouging. Zwolinski also deploys the nonworseness claim in his (2007) discussion of sweatshops, which is more clearly relevant to wage exploitation; but his comments in the price-gouging paper are, I think, more illuminating.

question-begging for someone to invoke this variant in response to a complaint that some employees don't consent to their employment. And it would be quite odd to claim that this variant is obviously true for consequentialists. It is only obviously true for consequentialists who deny the value of anything other than welfare—i.e., for welfarists. Likewise, the real nonworseness claim is only obviously true for consequentialists who deny the value of anything other than welfare and autonomy (or whatever undergirds the importance of consent). And there are plenty outside of that camp.

If the nonworseness claim isn't question-begging, something else must be going on here. To see what, it will be useful to consider our obligations as third parties—something the literature often focuses on (Wertheimer and Zwolinski 2015). Assuming we can't force Anne to hire Bob, the most we can do is ensure that *if* she hires him, she offers him better compensation—say, by instituting a minimum wage. Suppose we are confident that this would lead Anne to not hire Bob. Bob is already quite poorly off; surely we shouldn't do anything that will eliminate his best option, even to ensure that things are fair. Instead, it seems we should allow Anne to hire Bob, regardless of his compensation's fairness, so as to bring about more good.

If you are inclined to agree, this suggests that you share an evaluative intuition: the state of affairs in which Bob is consensually benefitted, even unfairly, is all-things-considered better than the state of affairs in which he is (fairly) left in poverty. From this, you might be tempted to infer a *deontic* claim: it cannot be *more wrong* for Anne to exploit Bob than for her not to hire him at all.

If we understand the motivations for the nonworseness claim this way, the passages above make much more sense. For Wertheimer, Anne is not begging the question, because she assumes her opponent will agree that it's better that Bob be treated unfairly than left in poverty, and would therefore endorse the relevant *evaluative* claim. And Wertheimer thinks Anne's opponent is likely to infer the relevant *deontic* claim, to accept that if A has “a right not to transact with B and . . . transacting

with B is not worse than not transacting with B, it can't be seriously wrong for [A] to engage in an unfair transaction with B."

Following this line, we can most charitably read Zwolinski as explaining why someone might make this inference: they're a consequentialist. For consequentialists, *more goodness means more rightness*, and thus it can't be wrong for Anne to bring about more overall value. On this reading, Zwolinski doesn't really mean that consequentialists have to endorse the nonworseness claim; he means that any consequentialist who accepts the relevant *evaluative* claim also has to accept the *deontic* one.

Since champions of the Reasonable View accept the evaluative claim, they must reject this inference. An easy way to do this, it might seem, would be to reject consequentialism. After all, for deontologists, it is not at all uncommon for more goodness *not* to entail more rightness. Consider Thomson's (1985) classic case of the surgeon who considers sacrificing one patient to save five others, often deployed as a counterexample to consequentialism. The deontologist's position is that killing the one to save the five is impermissible, despite the fact that it would generate an all-things-considered better outcome than (permissibly) allowing the five to die.

Yet, surprisingly, Zwolinski maintains that the nonworseness claim "is meant to have traction for deontological theories as well," because "it is puzzling how it could be worse by any moral standard . . . for *A* to provide some help than it is for him to provide none." Following Wertheimer, Zwolinski points out that to get around such puzzlement, deontologists need to defend an "interaction principle" according to which "there are special moral constraints that apply to our interactions with others *if we choose to interact with them*, but which do not have anything to say about those who choose not to interact at all" (Zwolinski 2008, 357).

Zwolinski finds such interaction principles implausible. Indeed, even Wertheimer, who defends them, proceeds with some trepidation. The reason, I suspect, is that it is tempting to think of

interaction principles as counterintuitively placing extra moral burdens on the virtuous, just as Zwolinski charges:

[T]heories which subject *A* to a moral obligation to rescue *B* on condition of their interaction seem to suffer from both a defect of unfairness and a serious internal tension. *A*'s interaction with *B* is supererogatory, done with *B*'s consent, violates no independent moral constraints, and benefits *B*. Why, then, should the interaction itself place *A* under new moral obligations toward *B*, beyond those to which *A* and *B* mutually agree? On the face of it, it seems unfair to burden *A* with this extra requirement given that he is *already* doing more than is morally required of him. (Zwolinski 2008, 359)

To make this concrete, consider an interaction principle that requires you to donate 10% of your income to charity, but only if you choose to donate anything at all. This does seem unfair; indeed, it seems downright bizarre. Presumably, something like the moral significance of welfare undergirds the fact (if it is a fact) that donating to charity is supererogatory. But if the moral significance of welfare merely *suggests* donating, why would it *require* donating more once you begin? This seems no less bizarre as a deontological side constraint than it would as an implication of a consequentialist calculus.⁹

That bizarreness, however, is not a feature of interaction principles per se; it is a feature of *monistic* interaction principles on which a single value requires more of you as you respond properly to it. Crucially, this is not the sort of interaction principle at work in the Reasonable View. The intuition behind the claim that Anne acts wrongly is not that Anne has to benefit Bob *more* because she has benefited him *some*. Rather, it is that she must benefit him more because she is keeping for herself an *unfair* portion of the profits created by their interaction. Anne's hiring Bob generates new obligations not because the moral significance of welfare requires more of her now that she's helped him, but because the moral significance of *fairness* has become salient within their interaction.¹⁰ Once we

⁹ Though for an attempt to make this more palatable under deontology, see Snyder (2008).

¹⁰ This is assuming that any prior unfairness placed no obligations on Anne.

understand this, the worry that interaction principles place unfair burdens on the generous vanishes. Anne runs afoul of her interaction obligation because she seeks to take too much for herself.

This is the crucial feature of the Reasonable View that the above criticisms miss. It is misleading to say that Anne is “providing help.” If that’s all she were doing, then *of course* it would be better than doing nothing. But her interaction with Bob is *mutually* beneficial, and fairness may well have something to say about how those benefits are to be distributed. This is simply a reflection of the fact that some values may generate new obligations within—perhaps only within—interactions.¹¹

Given this, we can vindicate the Reasonable View with the following, perfectly sensible deontological theory. Welfare is a good-making feature of states of affairs, but promoting the good is merely supererogatory, so Anne is not obligated to hire Bob. There are also side constraints: it is impermissible to distribute benefits unfairly, or to force people to do what they are morally permitted not to do (which is why we can’t make Anne hire Bob *and* compensate him fairly).

This possibility makes it clear that the intuitions used to motivate the nonworseness claim pose no principled threat to the Reasonable View, at least as a deontological theory. Again, this is not to say that the nonworseness claim is false, but only that it admits of no defense on grounds that are independent of any particular moral theory. If it is true, it is true because, say, only welfare and autonomy matter morally, or because they matter more than fairness. But that is precisely what champions of the Reasonable View can and should deny.

It might seem, though, that *only* deontologists can embrace the Reasonable View. Some, I expect, would find this a disappointing result. But Benjamin Ferguson takes an even dimmer view, arguing that the Reasonable View should be rejected because it faces the “paradox of deontology.” Recall that our deontological Reasonable View suggests that we are sometimes required to bring about

¹¹ “Any transaction or relationship that creates a social surplus gives rise to a new moral feature—fairness or unfairness—that does not arise outside of that transaction or relationship” (Wertheimer 1999, 291).

less of what's good. This is paradoxical, Ferguson charges, for "[p]rinciples that prohibit what they are designed to promote suffer from a form of inconsistency" (Ferguson 2011, 14).¹²

I suspect this complaint seems forceful here (to those who find it so) for the same reason the nonworseness claim does. Recall Thomson's surgeon case. The typical consequentialist response is not to call deontology paradoxical, but rather to resolve the apparent tension between the evaluative and the deontic, to either bite the bullet and insist that killing the one is right, or accept that killing the one is wrong but insist that the outcome of doing so wouldn't be better after all (e.g., by arguing that doctors' doing such things would discourage people from seeking medical attention). Relatedly, most would judge not only that it would be wrong for the surgeon to kill the one to save the five, but also that third parties should prevent the surgeon from doing so.

By contrast, for many the intuition that it would be better for Anne to exploit Bob than not to hire him survives reflection, as does the intuition that third parties should allow her to do so, even though it still seems Anne is doing something wrong. This may well exacerbate the appearance of tension between the Reasonable View's evaluative and deontic commitments. And this, in turn, may lend force to worries like Ferguson's, just as it lent force to the nonworseness claim.

Nevertheless, Ferguson's argument fails because it is question-begging. He seems absolutely right that the correct deontic principles should reflect what matters morally. But he simply assumes that *features of states of affairs* exhaust what matters morally. And that is just to assume that consequentialism is true. Suppose instead that *acting fairly* itself matters morally. It may be that this

¹² Ferguson's solution is to accept a more limited version of the nonworseness claim: "If one or more strictly Pareto improving and consensual transactions are possible, then for at least one of these transactions, transacting is morally better than not transacting" (Ferguson 2011, 15). I think this fails to sufficiently capture the intuitiveness of the evaluative aspect of the nonworseness claim, especially its plausibility as grounds for accepting that we must allow Anne to exploit Bob.

cannot be, or is not *best*,¹³ understood as a reflection of the value of some feature of states of affairs—e.g., the idea is not that we should maximize fair actions. That does not reveal some deep problem with theories that require acting fairly even when this would produce a worse state of affairs; it reveals only that they are not versions of consequentialism.¹⁴

¹³ This may all be moot if every moral theory can be “consequentialized.” Campbell Brown (2011) argues that they cannot. (He references Dreier (1993), Louise (2004) and Portmore (2007) as theorists who endorse the thesis that they can, “or something close to it” (Brown 2011, 750, fn.1).) However, Brown defines consequentialism in terms of “maximizing the good,” which is arguably too narrow. In consequentializing the Reasonable View in Part II, I take consequentialism to include any theory whose deontic commitments are entailed by its evaluative commitments, where value supervenes on features of (sets of) states of affairs. I make no argument here that all views can be consequentialized on this understanding of consequentialism. My suspicion is that they can, but that many—including, arguably, the Reasonable View—are not *best* understood as forms of consequentialism.

¹⁴ This is perhaps easiest to see on—but crucially does not *demand*—a buck-passing view on which facts about value are a function of facts about reasons, and (say) one always has decisive reason to distribute benefits fairly. Note that this is not just a problem for Ferguson, but for many (perhaps all) appeals to the “paradox of deontology.” For example, Alexander and Moore (2015) write that “[d]eontologists need their own, non-consequentialist model of rationality, one that is a viable alternative to the intuitively plausible, ‘act-to-produce-the-best-consequences’ model of [moral] rationality that motivates consequentialist theories. Until this is done, deontology will always be paradoxical.” But while deontologists do need a model of rationality, it is a mistake to claim that without one, their view is “paradoxical.” That is to question-beggingly judge deontological theories’ coherence against a background assumption that acting morally is about producing good outcomes. This is especially problematic given that deontologists are often working “bottom-up,” starting from *substantive* intuitions about what morality or rationality requires, ones the presumed consequentialist default can’t (easily) accommodate. All of this is symptomatic of the all-too-common, illegitimate assumption that to say that something matters morally (is valuable) is to say that it merits promotion as a feature of states of affairs. For a relevant, useful distinction between “Moorean” and “Kantian” value concepts see Bradley (2006).

Part II. Reasonable Consequentialism

In this second part of the paper, I demonstrate that even if one assumes that features of states of affairs exhaust what matters morally, one need not reject the Reasonable View. This exercise does not show that consequentialists *must* accept the Reasonable View, only that it is possible for them to do so. I develop a toy consequentialist theory, “Reasonable Consequentialism,” that entails the Reasonable View, given two further stipulations about Anne and Bob’s case. The first is that Anne’s exploiting Bob would decrease net fairness; the second is that forcing Anne to hire Bob would decrease net autonomy.¹⁵

Reasonable Consequentialism is consequentialist in that all of its deontic commitments are entailed by its evaluative ones, where value supervenes on sets of states of affairs.¹⁶ It is also, therefore, consistent with the idea that more goodness entails more rightness.

Because the Reasonable View involves supererogation—Anne isn’t required to hire Bob, though it would be good for her to do so—Reasonable Consequentialism is a form of *satisficing*

¹⁵ The first stipulation is fairly natural. It *could* turn out that in exploiting someone, we make things more fair overall. But given Bob’s poverty, it does not seem much of a stretch to suppose that the local unfairness of Anne’s exploitative offer would lead to a *globally* less fair outcome. The second stipulation might seem more problematic. But note that *any* consequentialist who accepts that we can’t force Anne to hire Bob *and* compensate him fairly, despite the increase in welfare this would arguably bring—including Zwolinski and other defenders of wage exploitation—will presumably have to vindicate a fairly general prohibition on forcing people to do supererogatory things (or take this prohibition to be merely contingent on the details of employment practices, which seems counterintuitive).

¹⁶ Some, of course, endorse narrower definitions of consequentialism. But note that it is already doubtful that we can rest easy with a traditional “maximize the good” understanding of consequentialism, since it is a near point of agreement in the relevant dialectic that benefitting Bob is supererogatory. In any case, whether or not Reasonable Consequentialism appeals to consequentialists, it avoids structural objections like Ferguson’s. Thanks to Daniel Nolan for discussion on this point.

consequentialism.¹⁷ For simplicity's sake, the satisficing level is set at the counterfactual status quo: actions that generate greater expected value than inaction are supererogatory; actions that generate less expected value than inaction are impermissible.¹⁸

Now, consider the following states of affairs:

EXPLOITED	The result of Anne's choosing to exploit Bob.
NO HIRE	The result of Anne's choosing not to hire Bob.
FAIR	The result of Anne's choosing to hire Bob and compensate him fairly.
FORCED	The result of Anne's being forced to hire Bob and compensate him fairly.

The view that we produce more good by allowing Anne to exploit Bob suggests that the gain in welfare is more important than the loss in fairness ($\text{EXPLOITED} > \text{NO HIRE}$). But if Anne is acting wrongly in exploiting Bob, then given consequentialism the result of her exploiting him must be worse than the result of her not hiring him ($\text{NO HIRE} > \text{EXPLOITED}$). That looks straightforwardly inconsistent. And that's why the Reasonable View seems to be incompatible with consequentialism.

But consequentialism is more flexible than this suggests. For it is perfectly possible for consequentialism to weight values differently in different contexts. Here, there is at least one relevant shift in context: the states of affairs accessible to us as third parties are different from those accessible to Anne. Everyone can access **EXPLOITED** and **NO HIRE**. Anne can make the relevant choices; third parties can incentivize her to make those choices. But only Anne can access **FAIR**, because (we are assuming) we have no way to incentivize Anne to willingly offer Bob fair compensation. And only we

¹⁷ See, e.g., Slote (1984).

¹⁸ Neither inaction nor negatively defined actions such as *not giving Bob a million dollars* count as actions here. If inaction were an action, determining the counterfactual status quo would require determining what action Anne would have taken if she took any action at all. If *not giving Bob a million dollars* were an action, Anne might be required to give Bob a million dollars, since not giving Bob a million dollars might well decrease welfare compared to the state of affairs resulting from her *not* not giving Bob a million dollars (depending how wealthy Anne is).

can access FORCED, because Anne can't force herself to do something. So, Anne can access {EXPLOITED, NO HIRE, FAIR}, while third parties can access {EXPLOITED, NO HIRE, FORCED}. What's more, there is a principled difference between these sets: only the latter involves *value conflict*—i.e., lacks a member that weakly dominates with respect to all values.

To see this, start with the states accessible to third parties: {EXPLOITED, NO HIRE, FORCED}. By doing nothing, third parties bring about EXPLOITED. Third parties can instead bring about NO HIRE by (say) instituting a minimum wage, thus incentivizing Anne not to hire Bob. Or, by forcing Anne to hire Bob, third parties can bring about FORCED. As seen in Figure 1, there is value conflict within this set: no member has at least as much welfare, fairness and autonomy as the others. Something must be sacrificed.

Figure 1. Expected value gains/losses for third parties*

	EXPLOITED [†]	NO HIRE	FORCED
Welfare	=	–	+
Fairness	=	+	=
Autonomy	=	=	–
*Shaded cells are the available maximum for a row (value). A shaded column (state of affairs) weakly dominates; here none do. †Gains/losses indexed to status quo (result of inaction)			

By contrast, consider the states accessible to Anne: {EXPLOITED, NO HIRE, FAIR}. By doing nothing, Anne remains in NO HIRE. By exploiting Bob, Anne can bring about EXPLOITED. Or by choosing to hire Bob and compensate him fairly, Anne can bring about FAIR. Here, there is no conflict, as seen in Figure 2. Anne can promote welfare without sacrificing anything of value, by choosing to hire Bob and compensate him fairly. FAIR weakly dominates.

Figure 2. Expected value gains/losses for Anne*

	EXPLOITED	NO HIRE [†]	FAIR
Welfare	+	=	+
Fairness	-	=	=
Autonomy	=	=	=
*Shaded cells are the available maximum for a row (value). A shaded column (state of affairs) weakly dominates; here FAIR does.			
†Gains/losses indexed to status quo (result of inaction)			

In order to entail the Reasonable View, Reasonable Consequentialism needs to have the following deontic implications.

- (1) It is permissible for Anne not to hire Bob.
- (2) It is impermissible for Anne to exploit Bob.
- (3) It is impermissible for third parties to raise the minimum wage (i.e., do anything expected to lead Anne not to hire Bob).
- (4) It is impermissible for third parties to force Anne to hire Bob.

Reasonable Consequentialism achieves this by calculating overall value with respect to sets of states of affairs,¹⁹ holding that (a) autonomy's²⁰ value is lexically prior to both welfare's and fairness'; (b) welfare's value is lexically prior to fairness' conditional on the *presence* of value conflict within the relevant set; and (c) fairness' value is lexically prior to welfare's conditional on the *absence* of value conflict within the relevant set.²¹

¹⁹ For deontic claims, the relevant set contains the states accessible to the agent in question, as seen below. For comparative evaluative claims, the relevant set contains the states being compared. Thus, for instance, because the set {EXPLOITED, NO HIRE} involves value conflict, Reasonable Consequentialism judges that EXPLOITED > NO HIRE with respect to this set, in keeping with the evaluative claim used to motivate the nonworseness claim.

²⁰ Or whatever underpins the impermissibility of forcing Anne to do what she is morally permitted not to do.

²¹ Holding autonomy fixed, this generates: Given a set of states of affairs {S₁-S_n}, S_X is all-things-considered better than S_Y iff (i) S_X contains more fairness but no less welfare than S_Y or (ii) S_X contains more welfare but

On Reasonable Consequentialism, it is always permissible to remain at the status quo, to take no action. Therefore, (1) it is permissible for Anne not to hire Bob.

If Anne exploits Bob, this would bring about greater welfare, but less fairness. Since there is no value conflict in the set of states accessible to Anne, the loss in fairness outweighs the gain in welfare, and so exploiting Bob brings about less overall value. Therefore, (2) it is impermissible for Anne to exploit Bob.

Third parties' raising the minimum wage is expected to lead Anne not to hire Bob, and would thus increase fairness but lower welfare. Since there is value conflict in the set of states accessible to third parties, the gain in welfare outweighs the loss in fairness, and so raising the minimum wage leads to lower expected value. Therefore, (3) it is impermissible for third parties to raise the minimum wage.

Finally, since autonomy always has priority, forcing Anne to hire Bob and compensate him fairly would also lead to lower expected value. Therefore, (4) it is impermissible for third parties to force Anne to hire Bob.²²

Of course, this conditional weighting scheme is ad hoc; but it is not unmotivated. It serves to capture, within a consequentialist framework, an intuition about how the value of fairness constrains us—an intuition typically associated with deontology. When one must choose between welfare and fairness, one should choose welfare. This is why, where welfare and fairness conflict, the former has priority. But when we are promoting welfare, we are required to distribute welfare gains as fairly as we

no less fairness than S_Y or (iii) S_X contains more welfare but less fairness than S_Y and there is *no* S_Z with no less welfare than S_X and no less fairness than S_Y or (iv) S_X contains more fairness but less welfare than S_Y and there is *some* S_Z with no less welfare than S_X and no less fairness than S_Y .

²² Reasonable Consequentialism's other implications are also in line with the Reasonable View. (5) By allowing Anne to exploit Bob, third parties create more welfare but less fairness. Given welfare's priority under value conflict, this means greater expected value, and is therefore supererogatory. (6) If Anne were to hire Bob and compensate him fairly, this would bring about greater welfare and have no negative impact on fairness or autonomy; it brings about greater expected value, and is therefore supererogatory.

can. This is why, where welfare and fairness do not conflict, it is impermissible to promote the former without maintaining the latter.

Reasonable Consequentialism entails and therefore serves as a possibility proof for consequentialist endorsement of the Reasonable View. It also avoids Ferguson's complaints. Because the view's deontic commitments are a direct consequence of its evaluative ones, it never recommends bringing about less of what it deems valuable; more goodness always means more rightness.

Conclusion

Defenders of wage exploitation offer powerful economic and moral arguments that the welfare- and autonomy-promoting power of apparently exploitative employment relations should make us reluctant to interfere with those relations.²³ Some of these theorists then draw on the nonworseness claim to defend the idea that not only should third parties refrain from interfering with these relations, but the employers themselves commit no serious wrong in entering into them. This *may* be true, but only insofar as it is defensible on substantive moral grounds—e.g, if welfare and autonomy are all that matters morally. Deontologists are welcome to reject these commitments, and thus to embrace the Reasonable View. Reasonable Consequentialism shows that consequentialists can do the same. This is an admittedly tortured view, suggesting that the Reasonable View is perhaps *best* understood as a form of deontology. But Reasonable Consequentialism nevertheless serves as further proof that there is no deep problem with the Reasonable View. It also highlights the potential moral significance of the fact that different states of affairs are accessible to different agents, even for consequentialists.

As discussed in the introduction, it is sometimes argued that in real-world cases of wage exploitation, employers don't get an unfair portion of the profits their employees generate. This too

²³ Though for an important response to some of the more popular defenses of wage exploitation, see Kates (2015).

may be true. But that is merely a contingent fact about some actual employment relations. Proponents of the nonworseness claim want us to accept the much stronger conclusion that it is not *possible* for cases like Anne and Bob's to be wrongfully exploitative. In this, they overreach. Perhaps they can convince us that we should not interfere with such relations because they are typically or even always more beneficial than harmful. But if so, they need not further convince us that they are not *wrongful*. They need only convince us that in some cases it is better to allow the wrong, to bring about the good.

Acknowledgements

Thanks to Michael Kates for work and discussions that inspired this paper. Thanks to Tristram McPherson and Daniel Wodak for helpful comments. And thanks to John Hasnas, Peter Jaworski, Govind Persad, Abe Singer and other participants in the Georgetown Institute for the Study of Markets and Ethics Faculty Workshop Series. Thanks also to the editors and three anonymous reviewers at *Business Ethics Quarterly*.

References

- Alexander, L., and M. Moore. 2015. "Deontological Ethics." In *The Stanford Encyclopedia of Philosophy*, edited by E. N. Zalta, Spring 2015.
<http://plato.stanford.edu/archives/spr2015/entries/ethics-deontological/>.
- Andrews, S. 1852. *The Science of Society*, No. 1. Fowlers & Wells.
- Arneson, R. J. 2013. "Exploitation and Outcome." *Politics, Philosophy & Economics* 12 (4): 392–412.
- Arnold, D. G., and N. E. Bowie. 2003. "Sweatshops and Respect for Persons." *Business Ethics Quarterly* 13 (2): 221–42.
- Bailey, A. 2010. "The Nonworseness Claim and the Moral Permissibility of Better-Than-Permissible Acts." *Philosophia* 39 (2): 237–50.
- Barnes, M. 2013. "Exploitation as a Path to Development: Sweatshop Labour, Micro-Unfairness, and the Non-Worseness Claim." *Ethics and Economics* 10 (1): 26–43.
- Bradley, B. 2006. "Two Concepts of Intrinsic Value." *Ethical Theory and Moral Practice* 9 (2): 111–30.
- Brown, C. 2011. "Consequentialize This." *Ethics* 121 (4): 749–71.
- Dreier, J. 1993. "Structures of Normative Theories." *Monist* 76 (1): 22–40.
- Ferguson, B. 2011. "The Paradox of Exploitation." *Erkenntnis*, 1–22.
- Kates, M. 2015. "The Ethics of Sweatshops and the Limits of Choice." *Business Ethics Quarterly* 25 (2): 191–212.

- Louise, J. 2004. "Relativity of Value and the Consequentialist Umbrella." *Philosophical Quarterly* 54 (217): 518–36.
- Malmqvist, E. 2015. "Better to Exploit than to Neglect? International Clinical Research and the Non-Worseness Claim." *Journal of Applied Philosophy*.
- Meyers, C. 2004. "Wrongful Beneficence: Exploitation and Third World Sweatshops." *Journal of Social Philosophy* 35 (3): 319–33.
- Portmore, D.W. 2007. "Consequentializing Moral Theories." *Pacific Philosophical Quarterly* 88 (1): 39–73.
- Powell, B., and M. Zwolinski. 2012. "The Ethical and Economic Case Against Sweatshop Labor: A Critical Assessment." *Journal of Business Ethics* 107 (4): 449–72.
- Sample, R. 2003. *Exploitation, What It Is and Why It Is Wrong*. Rowman & Littlefield.
- Slote, M. 1984. "Satisficing Consequentialism." *Proceedings of the Aristotelian Society* 58: 139–163.
- Snyder, J. 2010. "Exploitation and Sweatshop Labor: Perspectives and Issues." *Business Ethics Quarterly* 20 (2): 187–213.
- Snyder, J.C. 2008. "Needs Exploitation." *Ethical Theory and Moral Practice* 11 (4): 389–405.
- Thomson, J.J. 1985. "The Trolley Problem." *Yale Law Journal* 94 (6): 1395–1415.
- Valdman, M. 2009. "A Theory of Wrongful Exploitation." *Philosophers' Imprint* 9 (6): 1–14.
- Wertheimer, A. 1999. *Exploitation*. Revised edition. Princeton University Press.
- Wertheimer, A., and M. Zwolinski. 2015. "Exploitation." In *The Stanford Encyclopedia of Philosophy*, edited by E.N. Zalta, Summer 2015.
<http://plato.stanford.edu/archives/sum2015/entries/exploitation/>.
- Zwolinski, M. 2007. "Sweatshops, Choice, and Exploitation." *Business Ethics Quarterly* 17 (4): 689–727.
- . 2008. "The Ethics of Price Gouging." *Business Ethics Quarterly* 18 (3): 347–78.